

# Hashpipe Structures

Deepthi Gorthi

April 19, 2018

This file is a compilation of all the handy structures and functions Hashpipe provides for building a fast data capture and recording pipeline.

## 1 Thread Structures and Functions

### 1.1 hashpipe\_thread\_desc\_t

```
struct hashpipe_thread_desc {  
    const char * name;  
    const char * skey;  
    initfunc_t init;  
    runfunc_t run;  
    databuf_desc_t ibuf_desc;  
    databuf_desc_t obuf_desc;  
};
```

The `hashpipe_thread_desc` structure is used to store metadata describing a hashpipe thread. Typically a hashpipe plugin will define one of these hashpipe thread descriptors per hashpipe thread.

1. `const char *name`: String containing the thread name. Hashpipe threads are identified by their names which need to be registered with `register_hashpipe_thread()`. This is used to match command line thread specifiers to thread metadata so that the pipeline can be constructed as specified on the command line.
2. `const char *skey`: String containing the thread's status buffer "status" key. It is typically 8 characters or less, uppercase and ends with

“STAT”. If it is non-NULL and non-empty, HASHPIPE will automatically store/update this key in the status buffer with the thread’s status at initialization (“init”) and exit (“exit”).

3. `initfunc_t init`: Pointer to thread’s initialization function. The thread initialization function can be null if no special initialization is needed. If provided, it must point to a function with the following signature:

```
int my_thread_init_funtion(hashpipe_thread_args_t *args)
```

4. `runfunc_t run`: Pointer to thread’s run function. The thread run function must have the following signature:

```
void my_thread_run_funtion(hashpipe_thread_args_t *args)
```

5. `ibuf`: Structure describing thread’s input data buffer (if any)
6. `obuf`: Structure describing thread’s output data buffer (if any). The data buffer description structure used for `ibuf` and `obuf` currently contains one function pointer:

create - A pointer to a function that creates the data buffer.

Future HASHPIPE versions may introduce additional data buffer fields. `ibuf.create` should be NULL for input-only threads and `obuf.create` should be NULL for output-only threads. Having both `ibuf.create` and `obuf.create` set to NULL is invalid and the thread will not be used. The create function must have the following signature:

```
hashpipe_databuf_t *my_create_function(  
    int instance_id, int databuf_id)
```

## 1.2 hashpipe\_thread\_args

This structure passed (via a pointer) to the application’s thread initialization and run functions. The ‘user\_data’ field can be used to pass info from the init function to the run function.

1. `hashpipe_thread_desc_t *thread_desc`
2. `int instance_id`

3. `int` input\_buffer
4. `int` output\_buffer
5. `unsigned int` cpu\_mask: 0 means use inherited
6. `int` finished
7. `pthread_cond_t` finished\_c
8. `pthread_mutex_t` finished\_m
9. `hashpipe_status_t` st
10. `hashpipe_databuf_t` \*ibuf
11. `hashpipe_databuf_t` \*obuf
12. `void` \*user\_data

### 1.3 Useful Functions

1. `int` run\_threads(): Function threads used to determine whether to keep running.
2. `register_hashpipe_thread` (`hashpipe_thread_desc_t` \*ptm): Function should be used by pipeline plugins to register threads with the pipeline executable.
3. `hashpipe_thread_desc_t` \*find\_hashpipe\_thread(`char` \*name): This function can be used to find hashpipe threads by name. It is generally used only by the hashpipe executable. Returns a pointer to its `hashpipe_thread_desc_t` structure or NULL if a test with the given name is not found. Names are case sensitive.
4. `void` list\_hashpipe\_threads(FILE \*f): List all known hashpipe threads to file pointed to by the file pointer.
5. `unsigned int` get\_cpu\_affinity(): Get the CPU affinity of calling thread.

## 2 Data Buffer Structures and Functions

### 2.1 hashpipe\_databuf\_t

1. `char data_type[64]`: Type of data in the buffer
2. `size_t header_size`: Size of each block header in bytes
3. `size_t block_size`: Size of each data block in bytes.
4. `int n_block`: Number of data blocks in buffer
5. `int shmid`: ID of this shared memory segment
6. `int semid`: ID of locking semaphore set

### 2.2 Associated functions

1. **Key:**

`key_t hashpipe_databuf_key()`

Get the base key to use for **all** hashpipe databufs. The base key is obtained by calling the `ftok` function, using the value of `$HASHPIPE_KEYFILE`, if defined, or `$HOME` from the environment or, if `$HOME` is not defined, by using `tmp`. By default (i.e. no `HASHPIPE_KEYFILE` in the environment), this will create and connect to a user specific set of shared memory buffers (provided `$HOME` exists in the environment), but if desired users can connect to any other set of memory buffers by setting `HASHPIPE_KEYFILE` appropriately.

2. **Create Databuf:**

```
hashpipe_databuf_t *hashpipe_databuf_create(int
    instance_id, int databuf_id, size_t
    header_size, size_t block_size, int n_block)
```

Create a new shared mem area with given params. Returns pointer to the new area on success, or `NULL` on error. Returns error if an existing shmem area exists with the given `shmid` and different sizing parameters.

### 3. Get Databuf:

```
hashpipe_databuf_t *hashpipe_databuf_attach(int instance_id  
, int databuf_id)
```

Return a pointer to a existing shmем segment with given id. Returns error if segment does not exist

### 4. Detach Databuf:

```
int hashpipe_databuf_detach(hashpipe_databuf_t *d)
```

Detach from shared mem segment

### 5. Clear Databuf:

```
void hashpipe_databuf_clear(hashpipe_databuf_t *d)
```

### 6. Reset Pointer location:

```
char *hashpipe_databuf_data(  
    hashpipe_databuf_t *d, int block_id)
```

Returns pointer to the beginning of the given data block.

### 7. Get lock status:

```
int hashpipe_databuf_block_status(  
    hashpipe_databuf_t *d, int block_id)
```

```
int hashpipe_databuf_total_status(  
    hashpipe_databuf_t *d)
```

```
uint64_t hashpipe_databuf_total_mask(  
    hashpipe_databuf_t *d)
```

Returns lock status for given block\_id, or total for whole array.

### 8. Locking functions:

```
int hashpipe_databuf_wait_filled(  
    hashpipe_databuf_t *d, int block_id)
```

```
int hashpipe_databuf_busywait_filled(  
    hashpipe_databuf_t *d, int block_id)
```

```

int hashpipe_databuf_set_filled(
    hashpipe_databuf_t *d, int block_id)

int hashpipe_databuf_wait_free(
    hashpipe_databuf_t *d, int block_id)

int hashpipe_databuf_busywait_free(
    hashpipe_databuf_t *d, int block_id)

int hashpipe_databuf_set_free(
    hashpipe_databuf_t *d, int block_id)

```

Databuf locking functions. Each block in the buffer can be marked as free or filled. The "wait" functions block (i.e. sleep) until the specified state happens. The "busywait" functions busy-wait (i.e. do NOT sleep) until the specified state happens. The "set" functions put the buffer in the specified state, returning error if it is already in that state.

### 3 Status Handling

Routines dealing with the hashpipe status shared memory segment. Info is passed through this segment using a FITS-like keyword=value syntax.

```

typedef struct {
    int instance_id;
    int shmid;
    sem_t *lock;
    char *buf;
} hashpipe_status_t;

```

#### Status handling functions

1. Semaphore name: Stores the hashpipe status (POSIX) semaphore name in semid buffer of length size. Returns 0 (no error) if semaphore name fit in given size, returns 1 if semaphore name is truncated.

The hashpipe status semaphore name is `$HASHPIPE_STATUS_SEMNAME` (if defined in the environment) or `$HASHPIPE_KEYFILE_hashpipe_status`

(if defined in the environment) or \$HOME hashpipe\_status (if defined in the environment) or "/tmp/hashpipe\_status" (global fallback). Any slashes after the leading slash are converted to underscores.

```
int hashpipe_status_semname(int instance_id, char * semid,
size_t size)
```

2. Instance ID number:

Returns non-zero if the status buffer for instance\_id already exists.

```
int hashpipe_status_exists(int instance_id)
```

3. Attach buffer:

Return a pointer to the status shared mem area, creating it if it doesn't exist. Attaches/creates lock semaphore as well. Returns nonzero on error.

```
int hashpipe_status_attach(int instance_id, hashpipe_status_t
*s)
```

4. Detach buffer: Detach from shared mem segment

```
int hashpipe_status_detach(hashpipe_status_t *s)
```

5. Locking/Unlocking status buffer:

Lock/unlock the status buffer. hashpipe\_status\_lock() will sleep while waiting for the buffer to become unlocked. hashpipe\_status\_lock\_busywait will busy-wait while waiting for the buffer to become unlocked. Return non-zero on errors.

```
int hashpipe_status_lock(hashpipe_status_t *s)
```

```
int hashpipe_status_lock_busywait(
hashpipe_status_t *s)
```

```
int hashpipe_status_unlock(hashpipe_status_t *s)
```

6. Check formatting:

Check the buffer for appropriate formatting (existence of "END"). If not found, zero it out and add END.

```
void hashpipe_status_chkinit(hashpipe_status_t *s)
```

7. Clear: Clear out whole buffer

```
void hashpipe_status_clear(hashpipe_status_t *s)
```

8. Thread safe lock/unlock:

Thread-safe lock/unlock macros for status buffer used to ensure that the status buffer is not left in a locked state. Each `hashpipe_status_lock_safe` or `hashpipe_status_lock_busywait_safe` must be paired with a `hashpipe_status_unlock_safe` in the same function and at the same lexical nesting level. See "man `pthread_cleanup_push`" for more details.

**Note** This header file does not include `pthread.h` where `pthread_cleanup_push` and `pthread_cleanup_pop` are defined. Users of the macros defined here must explicitly include `pthread.h` themselves.

```
#define hashpipe_status_lock_safe(s) \
    pthread_cleanup_push((void (*)(void *)) \
        hashpipe_status_unlock, s); \
    hashpipe_status_lock(s);

#define hashpipe_status_lock_busywait_safe(s) \
    pthread_cleanup_push((void (*)(void *)) \
        hashpipe_status_unlock, s); \
    hashpipe_status_lock_busywait(s);

#define hashpipe_status_unlock_safe(s) \
    hashpipe_status_unlock(s); \
    pthread_cleanup_pop(0);
```



## 4 Socket and UDP packet handling

### 4.1 "hashpipe\_pktsock.h"

Routines for dealing with packet sockets. See `man 7 packet` and `packet_mmap.txt` for more info.

```
struct hashpipe_pktsock {
    unsigned int frame_size;
    unsigned int nframes;
    unsigned int nblocks;
    int fd;
    unsigned char *p_ring;
    int next_idx;
};
```

#### 4.1.1 Functional preprocessors

1. **Header:** Return header field 'h' from packet frame pointed to by 'p'  
`TPACKET_HDR(p,h) (((struct tpacket_hdr *)p)->h)`
2. **MAC:** Return pointer to MAC header inside packet frame pointed to by 'p'.  
`PKT_MAC(p) (p+TPACKET_HDR(p, tp_mac))`
3. **Network:** Return pointer to network (e.g. IP) packet inside packet frame pointed to by 'p'.  
`PKT_NET(p) (p+TPACKET_HDR(p, tp_net))`
4. **Check UDP:** Returns true (non-zero) if this is a UDP packet.  
`PKT_IS_UDP(p) ((PKT_NET(p)[0x09]) == IPPROTO_UDP)`
5. **Destination:** Returns UDP destination port of packet. **NB:** this assumes but does NOT verify that the packet is a UDP packet!  
`PKT_UDP_DST(p) (((PKT_NET(p)[0x16]) << 8) |  
 ((PKT_NET(p)[0x17]) >> 8))`

6. **UDP Size:** Returns size of UDP packet (including the 8 byte UDP header). **NB:** this assumes but does NOT verify that the packet is a UDP packet!

```
PKT_UDP_SIZE(p) (((PKT_NET(p)[0x18]) << 8) |
                 ((PKT_NET(p)[0x19])      ))
```

7. **Get UDP Payload:** Returns pointer to UDP packet payload. **NB:** this assumes but does NOT verify that the packet is a UDP packet!

```
PKT_UDP_DATA(p) (PKT_NET(p) + 0x1c)
```

#### 4.1.2 Useful Functions

1. Bind to Socket and attach it to a buffer:

`p_ps` should point to a `struct pktsock` that has been initialized by caller with desired values for the sizing parameters `frame_size`, `nframes`, and `nblocks`. `nblocks` MUST be a multiple of `nframes` and the resulting block size (i.e. `frame_size * nframes / nblocks`) MUST be a multiple of `PAGE_SIZE`.

`ifname` should specify the name of the interface to bind to (e.g. "eth2").

`ring_type` should be `PACKET_RX_RING` or `PACKET_TX_RING`.

Returns 0 for success, non-zero for failure. On failure, `errno` will be set.

Upon successful completion, `p_ps->fd` will be the file descriptor of the socket, and `p_ps->p_ring` will be a pointer to the start of the ring buffer.

```
int hashpipe_pktsock_open(struct
    hashpipe_pktsock *p_ps, const char *ifname,
    int ring_type)
```

2. Get frame:

Return pointer to frame or NULL if no frame ready.

If a non-NULL frame pointer is returned, the caller MUST release the frame back to the kernel (via `pktsock_release_frame`) once it is finished with the frame.

```

unsigned char *
    hashpipe_pktsock_recv_frame_nonblock(
struct hashpipe_pktsock *p_ps)

unsigned char * hashpipe_pktsock_recv_frame(
    struct hashpipe_pktsock *p_ps,
int timeout_ms)

```

3. Get matching frame:

If no frame is ready, returns NULL. If a non-matching frame is ready, it is released back to the kernel and NULL is returned. Otherwise, returns a pointer to the matching frame.

If a non-NULL frame pointer is returned, the caller MUST release the frame back to the kernel (via `pktsock_release_frame`) once it is finished with the frame.

```

unsigned char *
    hashpipe_pktsock_recv_udp_frame_nonblock(
struct hashpipe_pktsock *p_ps, int dst_port)

```

4. Wait for frame:

If no frame is ready, wait up to `timeout_ms` to get a frame. If no frame is ready after timeout, returns NULL. If frame does not match, releases non-matching frame back to the kernel and returns NULL. Otherwise returns pointer to matching frame.

If a non-NULL frame pointer is returned, the caller MUST release the frame back to the kernel (via `pktsock_release_frame`) once it is finished with the frame.

```

unsigned char * hashpipe_pktsock_recv_udp_frame(
    struct hashpipe_pktsock *p_ps, int dst_port,
int timeout_ms)

```

5. Release frame: Releases frame back to the kernel. The caller must be release each frame back to the kernel once the caller is done with the frame.

```

void hashpipe_pktsock_release_frame(unsigned char * frame)

```

6. Get stats: Stores packet counter and drop counter values in `*p_pkts` and `*p_drops`, provided they are non-NULL. This makes it possible to request one but not the other.

```
void hashpipe_pktsock_stats(struct
    hashpipe_pktsock *p_ps, unsigned int *p_pkts,
    unsigned int *p_drops)
```

7. Close socket: Unmaps kernel ring buffer and closes socket.

```
int hashpipe_pktsock_close(struct hashpipe_pktsock *p_ps)
```

## 4.2 hashpipe\_udp.h

```
struct hashpipe_udp_params {
    char sender[80];
    int port;
    char bindhost[80];
    int bindport;
    size_t packet_size;
    char packet_format[32];

    /* Derived from above: */
    int sock;
    struct addrinfo sender_addr;
    struct pollfd pfd;
};

struct hashpipe_udp_packet {
    size_t packet_size; /* packet size, bytes */
    char data[HASHPIPE_MAX_PACKET_SIZE]
    __attribute__((aligned(128)));
};
```

### Functions

1. Use sender and port fields in param struct to init the other values, bind socket, etc.

```
int hashpipe_udp_init(struct hashpipe_udp_params *p)
```

2. Close out socket, etc.

```
int hashpipe_udp_close(struct hashpipe_udp_params *p)
```

## 5 Error Handling

### 5.1 Exit codes

HASHPIPE_OK	0	Everything is great
HASHPIPE_TIMEOUT	1	Call timed out
HASHPIPE_ERR_GEN	-1	Super non-informative
HASHPIPE_ERR_SYS	-2	Failed system call
HASHPIPE_ERR_PARAM	-3	Parameter out of range
HASHPIPE_ERR_KEY	-4	Requested key doesn't exist
HASHPIPE_ERR_PACKET	-5	Unexpected packet size

### 5.2 Error logging

1. log error: Call this to log an error message.

```
void hashpipe_error(const char *name, const char *msg, ...)
```

2. log warning: Call this to log an warning message.

```
void hashpipe_warn(const char *name, const char *msg, ...)
```

3. log info: Call this to log an informational message.

```
void hashpipe_info(const char *name, const char *msg, ...)
```